

BACHELOR'S THESIS

**EXPLORING REASONING
PERFORMANCE OF RISC-V
SOFTWARE MODELS IN BTOR2**

by

NADIR FEJZIĆ

submitted in partial fulfillment of the requirements
for the degree of BACHELOR OF SCIENCE
in INFORMATICS

Department of Computer Sciences
Paris Lodron University of Salzburg
Salzburg, Austria

Supervised by:

Univ.-Prof. Dipl.-Inform. Dr.-Ing. Christoph Kirsch

August 16, 2024

Abstract

When working with software programs, we often want to ensure that they uphold certain invariants. There are various techniques for this, such as testing. However, testing is not exhaustive. Instead, we can reduce the problem of proving that the program upholds invariants to the boolean satisfiability problem. This reduction is done by encoding the program as a sequence of boolean formulas - known as the model, and solving the satisfiability of the resulting formula. Solving this problem is computationally expensive, as it is NP-Complete. In this thesis, we explore how different parameters of the generated models affect the performance of SAT solvers. In particular we present a benchmarking toolbox, and run benchmarks for different memory granularity parameters.

Contents

1	Introduction	4
2	Bounded model checking	5
2.1	Correctness of software programs	5
2.2	Reduction to SAT Problem	5
2.3	Complexity control with model parameters	7
3	Model generation	8
3.1	BTOR2 - brief introduction	8
3.2	rotor - tool of choice	10
3.2.1	Model Parameters and Checks	10
4	Experiment setup	11
4.1	peRIScope - short outline of available functionality	11
4.2	Workflow (peRIScope configuration, model generation, benchmarking, parsing of witness format, result files)	12
5	Experiment results	12
5.1	Results with binaries compiled with selfie	12
5.2	Results with binaries compiled with gcc	12
6	Conclusion and further work	12

1 Introduction

One of the requirements for software programs is the correctness. Varying techniques, such as testing, are used as an attempt to ensure the correctness of written software. However, in order to prove that the program is correct, one would have to test it with every possible input.

Doing that is not feasible for most of the practical programs. What we can do instead, is to prove that program has the desired properties for any input. Software programs are a series of machine instructions executed by a machine. Instructions and memory access in a machine are implemented using logical gates in hardware. These logical gates can be modelled as boolean formulas, which means that programs can be modelled as boolean formulas as well.

We can take advantage of this fact, and reduce the problem of proving the correctness of software programs to the boolean satisfiability problem. During the conversion of program to a boolean formula (the model) we can add constraints that model the desired properties of the program. Solving such model gives us the answer whether the program satisfies all of the constraints, and thus whether it possesses the desired properties.

Generating such models can be computed linearly in the size of the program (number of instructions). But since the model is a reduction of the correctness problem to a satisfiability problem, solving the model means solving the satisfiability problem. The satisfiability problem is, however, NP-Complete. This means that solving of such models is computationally expensive.

When generating models, we can tweak different parameters in order to influence its complexity. For example, machines contain circuits for memory access. Depending on the size of each addressable memory block (memory granularity), we need to access a different number of addresses. This means that, for smaller memory blocks, more addresses are needed to access the whole memory, which results in more complicated model of memory access. Larger memory granularity is better in this regard, but worse in others. In general, extracting values from memory that differ in number of bits from the used memory granularity increases complexity. We might need to use bit-shifting and bit-masking to extract the desired value. If our program has frequent access to values in memory smaller than the configured memory granularity, then this adds more complexity for each such memory access. So memory granularity might affect the solving performance.

We can tweak parameters of models to analyze how these shifts in complexity affect the performance of solvers. In this thesis we're particularly interested in solving performance when using different combinations of memory granularity for code and non-code memory segments. We also aim for a general setup that can be used to test and benchmark models of the same program generated using different parameters.

2 Bounded model checking

2.1 Correctness of software programs

Software programs are written with a particular goal in mind. We can create a specification that describes what the goal of the program is. If the program does what it's supposed to do as defined by the specification, we say that the program is correct. Depending on the program and our requirements, we can define different specifications for different types of correctness. For example, we can check whether the program logically does what it's supposed to. This would be logical correctness. Another example is we can check whether the program performs unsafe operations, such as accessing memory out of bounds. This would be safety correctness and so on.

In the context of this thesis, correctness of the program is defined as:

- Program does not terminate with a bad exit code
- Program terminates with a good exit code
- Program does not perform division by zero
- Program does not inhibit division overflow
- Program does not access any invalid memory address
- Program does not inhibit segmentation fault

However, testing for each of these properties for any input is problematic. If our program has a single 8-bit integer as its input, we would have to test it with 2^8 different inputs. For each bit added to the input, regardless if as an additional input or as an additional bit to the existing input, we double the number of tests. Testing with such large input space is not feasible for most of programs.

Better approach is to prove that our program is correct for any given input. In our case, each property we want to check for is a particular state we want our program to reach. What we're interested in, is whether an input for the given program exists, such that a machine reaches a certain state while it executes that program. This is known as the state reachability problem. In this thesis we further reduce the state reachability problem to the boolean satisfiability problem. The formal definition of the problem we want to solve is called model. The reduction is done in such a way, that our model is satisfiable if and only if the state is reachable. This means that, if our model is satisfiable, our program upholds the desired constraints.

2.2 Reduction to SAT Problem

A program is a series of instructions that can be executed by a particular machine. The machine decodes instructions and executes them. Instructions can have various semantics such as arithmetic operations, memory access, branching, and so on. By executing the instructions, machine changes its state. Minimal amount of state that is changed for each instruction is the update of the program counter. Apart from program counter, the machine can modify its registers and main memory as well.

Instructions in machines are implemented using logical gates, such as AND, OR, NOT and other gates. Logical gates are devices that perform boolean functions and are




Figure 1: Schematic of full adder implemented with two XOR gates, two AND gates, one OR gate.

implementations of boolean formulas. In other words, we can model the logical gates of the machine using boolean formulas. The logical circuit of one instruction becomes a boolean formula in our model. An example of a 1-bit full adder schematic can be seen in figure 1. The schematic shows circuit logic for arithmetic addition with two 1-bit inputs A and B , one 1-bit value from previous addition called carry in C_{in} and it computes the 1-bit sum S and 1-bit carry out C_{out} value in case of an overflow. If we have a simple RISC-V program containing the following series of instructions:

```
addi t0, zero, 4
addi t1, zero, 2
mul  t0, t0, t1
```

then we can model the program with the following operations:

Listing 1: Example model

```
t0_1 = 0 + 4 && pc_1 = pc_0 + 4 &&
t1_1 = 0 + 2 && pc_2 = pc_1 + 4 &&
t0_2 = t0_1 * t1_1 && pc_3 = pc_2 + 4
```

So the small RISC-V program can be reduced to the series of operations in listing 1. What we encoded is a satisfiability modulo theory formula (SMT-Formula). That is, we encoded the program as a series of operations in theory of bit-vectors. This model is satisfiable if and only if there exists an assignment of values to variables such that the formula is true. For the SMT-Formula bit-vectors and arrays of bit-vectors are used. Bit-vectors represent signed and unsigned integers of arbitrary bit-length, and bit vector arrays are arrays of N-bit bit-vectors as values indexable by an M-bit bit-vector. The model we encoded is an intermediate step, we need to further reduce it to boolean formula.

If we recall the full adder schematic in figure 1, we can see that the logic gates correspond to the following boolean formulas:

$$\begin{aligned} S &= A \oplus B \oplus C_{in} \\ C_{out} &= (A \wedge B) \vee (C_{in} \wedge (A \oplus B)) \end{aligned}$$

This is a series of boolean formulas that describe the behaviour of the 1-bit full adder. For two bit full adder, we would have two sum outputs S_1, S_2 , one initial carry in C_{in1} , and the carry out of the first sum C_{out1} would be the carry in to the second sum $C_{out1} = C_{in2}$. The same process applies for full adder of values with arbitrary bit length. The series of operations in listing 1 can be encoded in a similar way. This technique of encoding semantics of arithmetic and bitwise operations is called bit-blasting. We can add more constraints to the model that do not come from the program, but rather specify whether the machine reaches a particular state. For example, we can add constraints that check whether a certain register of the machine (variable) contains the value 0 and we perform a division with that register.

By doing so, we formulated a satisfiability problem, and by solving it we prove that either our program does not perform division by zero for any given input, or we get an example input for which the program performs division by zero. In particular, by doing this, we reduced the state reachability problem to the boolean satisfiability problem. Important observation is that the resulting boolean formula is finite and can model a finite number of executed instructions, which makes our models bounded in the number of executed instructions. This is the essence of bounded model checking.

2.3 Complexity control with model parameters

Model can be formulated for any given program with the technique of bit blasting. There are multiple possibilities when generating model where certain trade-offs can be made. In our models of the machine bit-vectors are used for values and arrays of bit-vectors for main memory and registers. As already mentioned, arrays contain bit-vectors of some fixed bit-length. If we want to model M -bit memory with an array that contains N -bit bit-vectors we need an address space of $2^{(M-N)}$ entries. More precisely, the array must be indexed with a bit-vector of bit length $(M - N)$. Size of bit-vectors stored in the array represent the memory granularity, the size of each memory block our machine can access. By changing the memory granularity we can directly impact the size of the array.

When accessing memory each memory block is addressable by an address, which is a bit-vector in our case. Logical circuit exists for this purpose in machine as well. Naively implemented, this could be a chain of something similar to if-else statements. In reality machines use a logical circuit called demultiplexer, which takes an input and routes it to one of several possible outputs according to an input binary address [1]. Demultiplexer is another logical circuit that has N inputs, each input having two possible values, and 2^N outputs. By translating the circuit logic of demultiplexer, we can model it as boolean formula as well, effectively modelling memory access of the machine. Arrays used for memory representation are indexed by N -bit bit-vectors, so the demultiplexer will have an N -bit input value. Complexity of the resulting circuit logic directly depends on the

size of the input, so choosing smaller memory granularity increases the complexity of the model as well.

On the other hand, if we choose larger memory granularity, we can access the whole memory with fewer addresses. Consequently, this reduces the memory access circuits and therefore the complexity in corresponding part of our model. However, it is not always the case that we need to access values in memory that have the same size as our chosen granularity. When that's not the case, we have to perform bit-shifting and bit-masking operations in order to load the correct value. For example, if we want to load a 32-bit value from memory with 8-bit granularity, we will have to access memory 4 times and perform bit-shifting to store parts of value at the correct position. Smaller granularity in this case results in four memory accesses through large circuit and four bit-shifting operations. If we, however, want to load a 32-bit value from memory with 64-bit granularity, we will have to perform a single load operation and bit-masking to trim the value to the correct size. Larger granularity in this case reduces number of times we need to access the memory and reduces the number of operations needed to correctly load the value.

In our models, we can control the memory granularity and therefore influence the complexity of models in various parts. In particular, we can control two different parts of memory: memory that holds the program instructions (code) and memory that holds the data (data, heap and stack segments). In particular, we try different combinations of code and data memory granularity.

3 Model generation

In order to properly model the machine, we need to generate a model that resembles it. Models need to precisely define the semantic of memory access and machine instructions. As we already mentioned, we do not generate boolean formula directly, but rather an SMT-Formula first. We use a special language for this purpose called BTOR2.

3.1 BTOR2 - brief introduction

BTOR2 is an extension of BTOR, which is a format for quantifier-free formulas over bit-vectors and arrays. BTOR2 extends BTOR by a set of additional features and includes witness, invariant and fairness constraints and liveness properties. BTOR2 supports definition of sorts, which can be thought of as types. For example, `sort bitvec 32` defines a sort of 32-bit bit-vectors. Each line in BTOR2 format starts with an integer, which is either a sort id for a sort definition, or node id for a node definition. Nodes in BTOR2 can define operations such as addition, multiplication etc., or they can be variable declaration, memory access, constraints and so on [4].

In listing 2 we can see an example of BTOR2 model. In first two lines we define two sorts, one is a 1-bit bit-vector with sort id 1, and the other is a 32-bit bit-vector with sort id 2. In the following line we have a node with node id 3 that defines an input variable `turn` of sort with sort id 1, meaning that `turn` is a 1-bit bit-vector. In the next line we define a constant 0 of sort with sort id 2. In next two lines we define two 32-bit bit-vector state variables `a` and `b`, and then initialize them with previously defined constant 0. We then define a constant 1, which we use to increment the state variables.

We can observe that most of the keywords first reference a sort id. The referenced sort id is the type of the node, and the node itself evaluates to a certain value that we can reference by the node's id. In nodes 12 and 13, we use an if-then-else construct, where we evaluate to a certain value based on the condition. For state variable **a** we check whether the value of node 3 is 1 (true), and if it is, the node 12 evaluates to value of node 5, otherwise it evaluates to the value of node 10. For state variable **b** we check whether the value of node 3 is false (-3 negates the value of node 3), and if so we use the value of node 6, otherwise value of node 11. Since nodes 5 and 6 are definitions of state variables, using them means using whatever the current value of the variable is. In nodes 14 and 15 we define the next function, which is the function that updates the state variables, and these functions use the previously defined if-then-else nodes for the update. In next few lines we define a 32-bit bit-vector constant 3, we have two nodes that check whether our state variables have value 3. If both of them have value 3, the node 19 evaluates to 1 (true), and in node 20 we define a bad state. The node 20 simply means, if the value of node 19 is 1, then this bad state is reached. In listing 3 we can see the corresponding C code for the BTOR2 model.

```

1 sort bitvec 1
2 sort bitvec 32
3 input 1 turn
4 zero 2
5 state 2 a
6 state 2 b
7 init 2 5 4
8 init 2 6 4
9 one 2
10 add 2 5 9
11 add 2 6 9
12 ite 2 3 5 10
13 ite 2 -3 6 11
14 next 2 5 12
15 next 2 6 13
16 constd 2 3
17 eq 1 5 16
18 eq 1 6 16
19 and 1 17 18
20 bad 19

```

Listing 2: Example BTOR2 model

```

int main(void) {
    bool turn;

    unsigned int a = 0;
    unsigned int b = 0;

    while (1) {
        turn = read_bool();

        assert(!(a == 3 && b == 3));

        if (turn)
            a += 1;
        else
            b += 1;
    }
}

```

Listing 3: Corresponding code for the model

The BTOR2 format encodes the SMT-Formula. To solve the formula, we use a set of tools that accompany the BTOR2 format contained in the project Boolector [3]. In particular, we use the reference implementation of a bounded model checker **btormc**. The bit-blasting of SMT-Formula is performed by the **btormc** before solving it. We run the model checker by providing it the file containing BTOR2 model, and optionally the upper bound. Upper bound option sets the maximal number of instructions that can be checked. If the model checker terminates with no output, that means that our program is does not satisfy the given constraints in the number of instructions we chose as the

upper bound. On the other hand, if the program satisfies the constraints, a witness format is produced. We can see an example output of witness format in listing ??.

```
sat
b0
@0
0 1 turn@0
@1
0 0 turn@1
@2
0 0 turn@2
@3
0 0 turn@3
@4
0 1 turn@4
@5
0 1 turn@5
@6
0 0 turn@6
.
```

Listing 4: BTOR2 Witness Format for model in listing 2

The **sat** indicates that the model is satisfiable, **b0** is the satisfiable property with **b0** for 0-th bad property, that is the first bad property that appears in the model. **@0** through **@6** indicate the assignments of the inputs in frames 0 to 6, e.g., e.g., in frame 0 the 0-th input (first input defined in the BTOR2 model) **turn** = 1 and so on. Frame 6 is the last frame, where the bad property **b0** is satisfied.

3.2 rotor - tool of choice

Generating BTOR2 models for programs is not done by hand. That would be tedious and error prone, as models tend to be very large counting tens of thousands of lines. Instead, we use tool capable of generating BTOR2 models. In particular, we use **rotor**, which is a self-translating modeling engine called based on **selfie** that translates full RISC-V code including all of **selfie** and itself to BTOR2 and SMT-LIB formulae that are satisfiable if and only if there is input to the code such that the code exits with non-zero exit codes, performs division by zero, or accesses memory outside of memory segments. Rotor also generates models that enable RISC-V code synthesis [2].

3.2.1 Model Parameters and Checks

rotor accepts multiple arguments and parameters that can be used to adjust produced models. By default multiple checks are enabled in **rotor**, and we have to explicitly disable them using following command line arguments:

- **-Pnobadexitcode** - Disables check for bad exit code. Without this option, the model checker will report satisfiable if the program can terminate with the chosen bad exit code.

- **-Pgoodexitcode** - Enables check for good exit code. With this enabled, the model is satisfiable if the program can exit with some exit code other than the chosen good exit code.
- **-Pnoexitcodes** - Disables checking whether multiple cores exit with the same exit code. If not present, model is satisfiable if different cores can exit with different exit codes.
- **-Pnodivisionbyzero** - Disables division by zero check, which checks whether an input for the program exists such that the program performs division by zero.
- **-Pnodivisionoverflow** - Disables check for division (or remainder) overflow. Without this option, model is satisfiable if an input exists such that program performs division with an overflow.
- **-Pnosegfaults** - Disables check for segfaults, which checks whether the program accesses memory outside of memory segments.

Apart from these checks, we can also modify some parameters of the model. Models generated for benchmarking in this thesis are generated with following parameters:

- **-bytestoread 1** - Input of the program is 1 byte
- **-cores 1** - Machine has a single core
- **-virtualaddressspace 32** - 32-bit virtual address space
- **-codewordsize X** - Granularity of the code memory segment, where X is the size of the code memory segment in bits
- **-memorywordsize X** - Granularity of the non-code memory segments (data, stack and heap), where X is the size of the memory segment in bits.
- **-heapallowance 4096** - Maximum heap size in bytes
- **-stackallowance 4096** - Maximum stack size in bytes

As we can see, the memory granularity of code and non-code memory segments can be independently configured. This was used to generate models with different combinations of code and non-code memory granularity. We can generate many models rather quickly, because the generation of models is linear to the size of the input program. **rotor** accepts either **C*** source code, binary compiled by **starc** or binary compiled by **gcc** as inputs. In case of using binaries as input, they must be compiled for the **RISC-V** architecture.

4 Experiment setup

4.1 peRIScope - short outline of available functionality

peRIScope is a benchmarking toolbox designed to make the process of parameterized model generation and benchmarking of model solving easier. The tool also supports parsing of BTOR2 witness format in order to provide more information, such as which property was satisfied, and transitions of inputs through frames that satisfy said property. **peRIScope** accepts a configuration file, which can be used to specify the parameters

for models, timeout for maximum duration of model solving, list of files for filtering, and flags for `btormc`. The `peRIScope` performs multiple run. In each run it reads the configuration, runs `rotor` with parameters of that run which results in multiple models being generated in (specific) `selfie` directory. Model generation is fast, so we can generate models for all example files with no filtering. After the models are generated, `btormc` is run on each model multiple times using `hyperfine` [5], measuring the time it takes each time. Example of the config file can be seen in listing 5.

```
timeout: 300
files: # perform model checking only on files with these names
  - "division-by-zero-3-35-rotorized.btor2"
  - "invalid-memory-access-fail-2-35-rotorized.btor2"

runs: # model configuration flags for 'rotor'
  8-bit-codeword-size: "0 -codewordsize 8"
  16-bit-codeword-size: "0 -codewordsize 16"
  32-bit-codeword-size: "0 -codewordsize 32"
  64-bit-codeword-size: "0 -codewordsize 64"
```

Listing 5: Example `peRIScope` configuration file

4.2 Workflow (`peRIScope` configuration, model generation, benchmarking, parsing of witness format, result files)

5 Experiment results

5.1 Results with binaries compiled with `selfie`

5.2 Results with binaries compiled with `gcc`

6 Conclusion and further work

References

- [1] Paul Horowitz, Winfield Hill, and Ian Robinson. *The art of electronics*. Vol. 2. Cambridge university press Cambridge, 1989.
- [2] Christoph Kirsch. *selfie*. August 2024. URL: <http://selfie.cs.uni-salzburg.at/>.
- [3] Aina Niemetz, Mathias Preiner, and Armin Biere. “Boolector 2.0”. In: *J. Satisf. Boolean Model. Comput.* 9.1 (2014), pp. 53–58. DOI: 10.3233/sat190101. URL: <https://doi.org/10.3233/sat190101>.
- [4] Aina Niemetz et al. “Btor2, btormc and boolector 3.0”. In: *International Conference on Computer Aided Verification*. Springer. 2018, pp. 587–595.
- [5] David Peter. *hyperfine*. Version 1.16.1. March 2023. URL: <https://github.com/sharkdp/hyperfine>.

Statement of Authentication

I hereby declare that I have written the present thesis independently, without assistance from external parties and without use of other resources than those indicated. The ideas taken directly or indirectly from external sources (including electronic sources) are duly acknowledged in the text. The material, either in full or in part, has not been previously submitted for grading at this or any other academic institution.

Salzburg, August 16, 2024

Nadir Fejzić